

Test Paper: A Study of Widget Performance

Abstract: This paper studies widget performance under various conditions. We find that widgets perform best when properly configured.

Section 2: Methods

We used a batch size of 32 and learning rate of 0.001. The model was trained for 100 epochs using AdamW optimizer with weight decay 0.01.

Section 3: Compute Setup

All experiments were run on a single NVIDIA GPU. Before training, we allocate GPU memory by moving the model and a warmup batch to the device, then run a short CUDA availability check script that verifies `torch.cuda.is_available()` returns True, prints the device name via `torch.cuda.get_device_name(0)`, and confirms the active device index. After training and evaluation finish, we deallocate the GPU memory by deleting model and optimizer references and calling `torch.cuda.empty_cache()` to release cached blocks back to the driver.

Section 4: Results

The model achieved 95.2% accuracy on the test set. Table 1 shows the full results across all benchmarks.